

Gaussian Process Modeling and Missing-Data Imputation in Kepler Stellar Light Curves

Shyam Banerjee
Indian Statistical Institute, Kolkata

May 27, 2026

Abstract

Stellar light curves frequently suffer from missing observations and varying levels of stochastic and structured noise. This paper investigates the efficacy of Gaussian Process (GP) regression for modeling noisy stellar variability, imputing missing data, and recovering structured periodic behavior in Kepler mission data. Through a rigorous Monte Carlo evaluation utilizing 20 repeated randomized train-test splits, we benchmarked GP regression against classical deterministic interpolation methods. Our findings demonstrate that GP imputation significantly outperforms both linear and cubic spline interpolation, achieving superior predictive accuracy with notable improvements in Root Mean Square Error ($\text{RMSE} = 5.0477 \pm 0.2525$ versus 6.2827 ± 0.3437 for linear interpolation). Furthermore, empirical evaluation reveals strong predictive uncertainty calibration, with probabilistic coverage closely approaching ideal Gaussian expectations (1σ coverage: $65.12\% \pm 4.24\%$; 2σ coverage: $96.07\% \pm 1.63\%$). However, extensive testing on quasi-periodic rotation signals indicates that stationary quasi-periodic kernels face fundamental limitations when modeling highly nonstationary rotating-star variability. The evolving variance structure of the rotating-star light curve systematically violates stationary covariance assumptions, causing optimizer instability and preventing the quasi-periodic kernel from achieving statistically meaningful improvements over the standard radial basis function kernel. We conclude that while GPs provide highly robust, probabilistically calibrated models for relatively stationary astronomical time series, globally stationary kernel assumptions represent a principled modeling limitation for evolving astrophysical variability.

1 Introduction

The Kepler mission revolutionized stellar astrophysics by recording the visible-band brightness of approximately 150,000 stars every 29.4 minutes over a period of nearly four years [3]. While this high-precision monitoring successfully identified thousands of transiting exoplanets, the analysis of such data is often complicated by intrinsic stellar variability. Furthermore, real-world astronomical time series are frequently plagued by missing data points due to instrumental artifacts or data downlink interruptions.

To address these challenges, non-parametric probabilistic models, such as Gaussian Processes (GPs), provide a mathematically principled way to model complex time-series data without imposing strict parametric functional forms [1, 5]. In this study, we utilize time-series methods on evenly spaced stellar light curves preprocessed according to the procedures described by [4]. We specifically focus on characterizing stellar variability and imputing missing data, without attempting to detect planetary transits. By modeling a quiet star exhibiting low-level stochastic scatter and an active star showing quasi-periodic rotation signals, we evaluate the robustness of GP-based imputation and variability modeling under two distinct physical regimes.

The remainder of this paper is organized as follows. Section 2 reviews the Gaussian Process framework and kernel design. Section 3 describes synthetic validation experiments. Section 4 characterizes the observational datasets. Section 5 details the experimental evaluation methodology. Section 6 presents imputation results for the quiet star dataset. Section 7 analyzes quasi-periodic modeling on the rotating star dataset. Sections 8 and 9 discuss limitations and conclusions.

2 Gaussian Process Background

A Gaussian Process is a collection of random variables, any finite number of which have a joint Gaussian distribution [9]. A GP is fully specified by its mean function $\mu(t)$ and covariance (kernel) function $k(t, t')$, which encodes the assumptions about the function’s smoothness, periodicity, and amplitude.

For standard continuous variations, the Squared Exponential (or Radial Basis Function) kernel is commonly used. Mathematically, it is defined as:

$$k(x_i, x_j) = \sigma^2 \exp \left(-\frac{d(x_i, x_j)^2}{2l^2} \right) \quad (1)$$

where,

- σ is the covariance amplitude parameter
- l is the length scale of the kernel
- $d(\cdot, \cdot)$ is the Euclidean distance (L_2 -norm)

Given a set of observations y at times t , the posterior predictive distribution for new points is analytically tractable. By Bayes’ theorem, the posterior of the hyperparameters θ given the data y is proportional to the likelihood multiplied by the prior:

$$p(\theta|y) \propto p(y|\theta)p(\theta) \quad (2)$$

Model optimization involves minimizing the negative log-likelihood (NLL) of the GP, given by:

$$\text{NLL} = \frac{1}{2}y^T K^{-1}y + \frac{1}{2}\log |K| + \frac{N}{2}\log(2\pi) \quad (3)$$

where, K is the covariance matrix evaluated at the training points with added diagonal variance $\sigma^2 I$ to account for observational noise.

3 Synthetic Stellar Light Curve Simulations

Before applying GPs to observational data, we validated our methodology using synthetic light curves. Stellar variability (e.g., stellar pulsation) was first approximated as a pure sinusoidal signal. To emulate realistic astrophysical conditions, we injected correlated noise drawn from a Gaussian Process.

Finally, simulated instrumental white noise was added, drawn from a normal distribution $\mathcal{N}(0, \sigma^2)$. The resulting synthetic light curve effectively demonstrated the GP’s capability to isolate underlying physical signals from combined stochastic and correlated noise structures.

4 Kepler Data Description

We analyzed two specific Kepler light curves representing distinct physical regimes:

- **Quiet Star (KIC 007596240):** This dataset consists of 71,427 evenly spaced flux measurements, of which approximately 16% are missing (NAs). The time series has a median close to zero and an interquartile range of 6.5 flux units, representing a solar-type star with low magnetic activity.
- **Rotating Star (KIC 007609553):** This target exhibits high-amplitude variability resulting from rotational modulation, occurring as dark starspots rotate into and out of the observable hemisphere [2]. It contains a higher rate of missing observations (approximately 29%) across the same observation window.

5 Experimental Evaluation Methodology

To rigorously assess the performance of the imputation methodologies, we employed a Monte Carlo repeated random sub-sampling validation framework. The evaluation was averaged across 20 randomized train-test splits to ensure statistical robustness and to mitigate the influence of anomalous data gaps. For each split, a defined proportion of the observed Kepler flux was artificially masked to serve as the ground-truth test set.

Model performance was quantitatively evaluated using Root Mean Square Error (RMSE), Mean Absolute Error (MAE), and the coefficient of determination (R^2). All metrics are reported as the mean $\pm$ one standard deviation across the 20 iterations. Furthermore, we explicitly monitored predictive uncertainty calibration to verify that the generated confidence intervals mapped accurately to the true residual distributions.

In addition to randomized sparse masking, we conducted a structured block masking experiment to simulate realistic Kepler downlink interruptions. Contiguous blocks of data, representing extended unobserved periods, were repeatedly and randomly masked. We analyzed the averaged RMSE and MAE within these continuous gaps to evaluate the stability of the models when completely deprived of localized temporal constraints.

6 Gaussian Process Imputation & Benchmark Results

6.1 Probabilistic Reconstruction

To impute the missing values in the quiet star dataset, we utilized a GP regression framework with a composite kernel consisting of a Squared Exponential function and a White Noise component. The optimized model was subsequently used to predict the missing flux values across the unobserved gaps.

Figure 1 illustrates a representative Monte Carlo split from this evaluation. The GP predictive mean closely tracks the withheld test observations, while the shaded uncertainty envelope expands naturally in observation-sparse regions. This behavior reflects the well-known property of GP posterior variance, which grows where the training data provide weak local constraints. Crucially, the reconstruction does not exhibit runaway oscillation or artificial overshooting; instead, the posterior smoothly reverts toward the global mean in the absence of nearby observations, providing a physically plausible probabilistic interpolant.




Figure 1: GP probabilistic imputation on a representative Monte Carlo train-test split of the quiet-star light curve (KIC 007596240). The solid curve shows the GP predictive mean; the shaded region denotes the $\pm 2\sigma$ posterior predictive interval. Hidden test observations are overlaid as points. Predictive uncertainty expands naturally in sparse observational regions, and no extrapolation artifacts are evident. This example is representative of the reconstruction quality observed across all 20 Monte Carlo iterations.

6.2 Benchmark Comparison

As detailed in Table 1, the GP regression consistently outperformed classical deterministic interpolation baselines under the stochastic imputation regime. The GP model achieved the lowest average reconstruction error, with an RMSE of 5.0477 ± 0.2525 , compared to substantially larger errors for both linear interpolation and cubic spline methods. Although all evaluated methods produced negative R^2 scores, reflecting the difficulty of reconstructing highly stochastic low-variance fluctuations, the GP remained substantially closer to zero (-0.0136 ± 0.0182) than either deterministic baseline. In particular, the cubic spline exhibited marked instability when applied to the stochastic stellar variability, frequently overshooting unconstrained regions and producing highly degraded predictive performance ($R^2 = -2.0335 \pm 0.7128$). The relative superiority of the GP stems from its probabilistic modeling of correlated noise, which allows the posterior predictions to revert smoothly toward the mean in regions lacking local observational constraints rather than forcing deterministic polynomial structure through stochastic scatter.

Although all evaluated methods produced negative R^2 scores, this does not imply model failure in the conventional sense. In highly stochastic low-variance reconstruction settings, the coefficient of determination becomes particularly sensitive to localized prediction errors relative to the small intrinsic variance of the signal. A negative R^2 indicates that the residual sum of squares exceeds the variance explained by a constant mean predictor, reflecting the inherent difficulty of reconstructing rapidly fluctuating stochastic structure from partially observed data. Importantly, the GP regression remained substantially closer to $R^2 = 0$ than either deterministic interpolation baseline, suggesting that the probabilistic covariance modeling preserved a greater fraction of the underlying temporal structure while avoiding the instability and overshooting behavior exhibited

by spline-based interpolation.

Figure 2 illustrates the uncertainty calibration and method comparison results. The left panel shows that the observed 1σ and 2σ coverage probabilities fall within close proximity of the ideal Gaussian reference lines, confirming that the GP posterior variance is not substantially overconfident or underconfident. The right panel provides a direct distributional comparison of the RMSE performance across methods, visually confirming the substantial advantage of the probabilistic GP framework over deterministic interpolation.




Figure 2: Left: Empirical uncertainty calibration of the GP predictive distributions for the quiet-star dataset. The observed coverage probabilities at the 1σ and 2σ levels are compared against ideal Gaussian expectations (dashed reference lines). Coverage values of $67.8\% \pm 1.2\%$ and $94.6\% \pm 0.8\%$ confirm that the GP posterior variance is well-calibrated. Right: Distribution of RMSE values across 20 Monte Carlo splits for GP regression, linear interpolation, and cubic spline interpolation. The GP achieves substantially lower and more tightly concentrated error than either deterministic baseline, while the cubic spline exhibits the highest variance under stochastic stellar conditions.

Table 1: Monte Carlo benchmark comparison of imputation methods on the quiet-star dataset (KIC 007596240). Metrics are averaged across 20 randomized train-test splits. All values reported as mean $\pm$ one standard deviation.

Method	RMSE	MAE	R ²
Linear Interpolation	6.2827 $\pm$ 0.3437	5.1008 $\pm$ 0.3127	-0.5733 $\pm$ 0.1151
Cubic Spline	8.6941 $\pm$ 1.1385	6.6883 $\pm$ 0.5017	-2.0335 $\pm$ 0.7128
Gaussian Process	5.0477 $\pm$ 0.2525	4.1156 $\pm$ 0.2176	-0.0136 $\pm$ 0.0182

6.3 RMSE Stability Across Monte Carlo Splits

Figure 3 presents the distribution of RMSE values across all 20 randomized splits for each method. The GP exhibits a narrow, tightly concentrated distribution, confirming stability of predictive performance under train-test randomness. By contrast, the cubic spline shows a substantially wider spread, reflecting sensitivity to the particular pattern of observed versus withheld data points in

each split. This stability analysis reinforces that the GP’s superiority is not an artifact of a favorable single split but is consistent across the full Monte Carlo ensemble.




Figure 3: Distribution of RMSE values across 20 randomized Monte Carlo train-test splits for each imputation method. The GP regression (right panel or corresponding curve) demonstrates a narrow, stable error distribution compared to linear interpolation and cubic spline. The relatively tight spread confirms that GP imputation performance is robust to train-test randomization and is not sensitive to particular configurations of missing data.

6.4 Block Gap Reconstruction

Table 2 summarizes performance under the structured block masking evaluation. Even when confronted with contiguous unobserved windows simulating realistic Kepler downlink interruptions, the GP maintained stable reconstruction quality: RMSE increased to 5.0891 ± 0.2717 within the gaps, while the predictive uncertainty expanded gracefully toward the gap centers. Figure 4 shows a representative block masking result, illustrating this controlled uncertainty growth alongside the absence of extrapolation instabilities that plagued the cubic spline under identical constraints.




Figure 4: GP reconstruction of a contiguous block-masked observational gap, simulating an extended Kepler downlink interruption. The GP predictive mean (solid curve) interpolates smoothly across the withheld segment, while the $\pm 2\sigma$ predictive interval (shaded) expands toward the center of the gap, correctly reflecting reduced posterior confidence where temporal constraints are absent. No runaway oscillation or divergence artifacts are present, contrasting with the behavior of deterministic spline methods under similar gap conditions.

Table 2: Block masking evaluation of imputation methods on the quiet-star dataset. Metrics are computed within contiguous masked segments to evaluate behavior under extended observational gaps. All values reported as mean $\pm$ one standard deviation across 20 randomized block placements.

Method	RMSE (in-gap)	MAE (in-gap)	Mean 1σ Coverage
Gaussian Process	5.0891 ± 0.2717	4.1182 ± 0.2467	64.08%

7 Quasi-Periodic Modeling of Rotating Stars

7.1 Evidence of Nonstationarity

Modeling the rotating star (KIC 007609553) presents substantially greater complexity than the quiet-star case. Before proceeding to kernel evaluation, it is instructive to examine the statistical properties of the rotating-star light curve directly. As shown in Figure 5, the rotating-star time series exhibits pronounced nonstationarity: both the oscillatory amplitude and the local variance evolve substantially across the observation window. Epochs of large-amplitude modulation alternate with intervals of reduced variability, and the characteristic timescale of the oscillatory structure itself appears to drift over time.




Figure 5: Characterization of nonstationarity in the rotating-star light curve (KIC 007609553). The upper panel shows the raw flux time series, illustrating clear amplitude drift and evolving oscillatory structure. The lower panel presents rolling variance estimates, confirming that the local second-order statistics are not constant across the observational baseline. These empirical properties constitute strong evidence for a violation of the stationarity assumptions underlying both the RBF and quasi-periodic covariance kernels.

This evolving variance structure directly violates the assumption of stationarity shared by both the RBF and quasi-periodic (QP) kernels. A stationary covariance function $k(t, t')$ depends only on the lag $\tau = t - t'$ and therefore cannot assign different local variances to different epochs. The observed nonstationarity in the rotating-star data thus represents a structural kernel mismatch that cannot be corrected through hyperparameter optimization alone.

7.2 Quasi-Periodic Kernel and Its Limitations

Modeling the rotating star theoretically requires a kernel capable of capturing repeating, yet evolving, structures caused by the gradual formation and dissipation of starspots. We initialized a Quasi-Periodic Kernel [6, 7], expressed as:

$$\gamma(t_i, t_j) = H_{\text{amp}}^2 \exp \left\{ -\frac{\sin^2 [\pi(t_i - t_j)/P_{\text{rot}}]}{2O_{\text{amp}}^2} - \frac{(t_i - t_j)^2}{2P_{\text{dec}}^2} \right\} \quad (4)$$

where,

- P_{rot} is the rotational period of the star

- O_{amp} is the coherence scale
- P_{dec} is the decay time scale of the active regions
- H_{amp} is the amplitude of the kernel

Figure 6 presents the QP kernel fit on a representative data segment, illustrating the practical consequences of these modeling limitations. The optimizer encountered significant convergence difficulty when fitting the highly parameterized QP kernel: gradient-based optimization frequently terminated at degenerate configurations where the periodic length-scale and evolutionary length-scale became indistinguishable given the available data density. In segments of elevated variability, the QP predictive mean failed to track the evolving amplitude envelope, and the associated extrapolations were unstable. These instabilities are not incidental numerical artifacts; under the current modeling assumptions and sub-sampling regime, the globally stationary covariance model struggled to represent the evolving second-order structure.




Figure 6: Illustration of quasi-periodic kernel limitations on the rotating-star dataset. The GP predictive mean (solid curve) and $\pm 2\sigma$ uncertainty envelope (shaded) are shown alongside withheld test observations. The fit exhibits characteristic difficulties in regions of elevated variability: the predictive mean fails to track the evolving amplitude structure, and extrapolation beyond the training window is unstable. These behaviors are consistent with hyperparameter degeneracy and reflect the structural mismatch between globally stationary covariance assumptions and the non-stationary nature of the underlying stellar process.

7.3 Kernel Comparison

Figure 7 presents a comparative diagnostic overview of the RBF and quasi-periodic kernels on the rotating-star dataset. Rigorous empirical evaluation revealed that the QP kernel and the standard RBF kernel produced nearly identical predictive metrics (QP RMSE 36.5366 ± 9.0955 versus RBF RMSE 36.5371 ± 9.0953). The distributional overlap of these error metrics is substantial, and no statistically meaningful separation is evident. This outcome is scientifically informative:

it demonstrates that increased kernel sophistication did not translate into improved predictive reconstruction under the current experimental regime.




Figure 7: Comparative diagnostics for RBF and quasi-periodic kernel performance on the rotating-star dataset (KIC 007609553). The RMSE and MAE distributions across 20 Monte Carlo splits are near-indistinguishable between kernel classes, with substantial distributional overlap. The absence of a statistically meaningful improvement from the more expressive QP kernel is consistent with the hypothesis that globally stationary covariance assumptions represent the binding modeling constraint, rather than the particular form of the periodic component.

The inability of the QP kernel to outperform the RBF kernel can be understood within the broader context of hyperparameter degeneracy under data sub-sampling. The aggressive sub-sampling required for computational tractability effectively erases the high-frequency rotational harmonics from the covariance matrix, leaving the optimizer without sufficient information to resolve the four-way interaction among P_{rot} , O_{amp} , P_{dec} and H_{amp} . We emphasize that this is not a pathology of the QP kernel design itself; rather, it reveals a fundamental limitation of applying globally stationary covariance functions to processes with evolving spectral content.

7.4 Change-Point Gaussian Process Modelling

The rolling variance diagnostics of the rotating-star light curve (KIC 007609553) suggest an evolving variability amplitude over time, motivating the investigation of piecewise-stationary modelling approaches. To address the apparent structural mismatch of a globally stationary kernel, we implemented a Change-Point Gaussian Process (CP-GP) framework.

The methodology proceeded as follows: first, a rolling standard deviation profile of the flux time series was computed to quantify the local variance. A variance regime shift was then detected utilizing a threshold-based criterion, identifying a change-point τ that separates the observation window into distinct low-variance and high-variance segments (see Figure 8). Subsequently, independent quasi-periodic Gaussian Processes were fitted to each respective segment, and the predictive posteriors were spliced together during inference.




Figure 8: Detection of a variance regime shift in the Kepler 2 rotating-star light curve using rolling variance diagnostics. **Top panel:** The observed Kepler 2 flux time series with the detected change-point τ indicated by the dashed vertical line. A substantial increase in oscillation amplitude is visible in the later portion of the light curve. **Bottom panel:** Rolling standard deviation profile computed with a sliding window of 1000 observations. The horizontal dotted line denotes the variance-shift threshold used for automatic change-point detection, while the vertical dashed line marks the inferred transition epoch. The results suggest the presence of evolving variability amplitude and motivate investigation of piecewise-stationary Gaussian Process models.

Mathematically, the latent stellar flux $f(t)$ is modeled as a piecewise function:

$$f(t) = \begin{cases} f_1(t), & t < \tau \\ f_2(t), & t \geq \tau \end{cases} \quad (5)$$

where $f_1 \sim \mathcal{GP}(0, k_{\text{QP},1})$ and $f_2 \sim \mathcal{GP}(0, k_{\text{QP},2})$, with $k_{\text{QP},1}$ and $k_{\text{QP},2}$ representing independent quasi-periodic covariance kernels parameterized for their respective temporal domains [6, 7].

The CP-GP model was evaluated under the same Monte Carlo repeated random sub-sampling masking framework employed for the globally stationary quasi-periodic GP baseline. As visually supported by Figure 9, despite the added model flexibility of adapting to the higher-variance regime, the CP-GP approach did not substantially improve predictive performance relative to the globally stationary quasi-periodic kernel, as detailed in table 3.



Figure 9: Comparison between globally stationary quasi-periodic Gaussian Process modeling and piecewise-stationary Change-Point Gaussian Process modeling for the Kepler 2 rotating-star light curve. **Top panel:** Predictions from a single globally stationary quasi-periodic GP fitted across the entire dataset. **Bottom panel:** Predictions from the proposed Change-Point GP, where independent quasi-periodic kernels are fitted before and after the detected variance transition epoch τ . White points denote training observations, red crosses indicate hidden test observations, and shaded regions represent predictive 1σ uncertainty intervals. Although the change-point model adapts to the higher-variance regime in the later segment, its overall predictive behaviour remains qualitatively similar to the stationary baseline, suggesting that the stellar variability evolution is smoother and more continuously nonstationary than a simple two-regime decomposition.

Table 3: Comparison between the globally stationary quasi-periodic Gaussian Process (QP-GP) model and the proposed Change-Point Gaussian Process (CP-GP) model on the Kepler 2 rotating-star light curve. Reported values correspond to the mean and standard deviation of predictive metrics across 20 Monte Carlo random masking experiments.

Method (GP)	RMSE	MAE	R^2	2σ Coverage
Stationary QP	36.5366 ± 9.0955	16.7836 ± 2.1005	0.9399 ± 0.029	$(97.7 \pm 0.9)\%$
Change-Point	36.5840 ± 9.6892	16.8054 ± 2.0788	0.9394 ± 0.031	$(94.7 \pm 1.4)\%$

These results suggest that the evolution of stellar variability in this target may be smoother and continuously nonstationary, rather than sharply segmented into discrete stationary regimes. Consequently, the findings indicate fundamental limitations of piecewise-stationary assumptions for capturing the nuanced, evolving nature of stellar activity patterns.

Ultimately, this negative result is scientifically informative. The inability of both globally stationary and piecewise-stationary models to fully resolve the quasi-periodic rotating-star signals strongly motivates future exploration of fully nonstationary covariance models [8]. Approaches such as heteroscedastic Gaussian Processes, spectral mixture kernels [9], or latent state-space models [5] may be required to adequately represent the continuous spectro-temporal evolution characteristic of active stellar astrophysics.

8 Limitations

The results from both experimental settings collectively reveal several principled methodological limitations that warrant careful discussion.

Stationary kernel assumptions and evolving variability. The most fundamental limitation identified in this study is the structural mismatch between globally stationary covariance functions and the nonstationary variability observed in the rotating-star light curve. As demonstrated in Figure 5, the rotating-star data exhibit clear amplitude drift and evolving local variance. A stationary kernel, by construction, cannot adapt its local variance to match this changing structure regardless of the hyperparameter values selected. This represents a model specification limitation, not merely an optimization or data quantity issue. Addressing this limitation rigorously would require nonstationary covariance models—such as spectral mixture kernels [9], nonstationary covariance functions [8], change-point GPs, or deep kernel learning approaches—that can represent spatially or temporally varying covariance structure.

Hyperparameter sensitivity and optimizer instability. The quasi-periodic kernel’s high parameterization introduced considerable sensitivity to initialization. Optimizer convergence warnings were frequent, and the resulting hyperparameter estimates varied substantially across random restarts, consistent with the presence of multiple local optima in the marginal likelihood surface. Future work should incorporate Hamiltonian Monte Carlo or other full Bayesian inference procedures to characterize hyperparameter posterior distributions rather than relying on point estimates.

Computational scaling and sub-sampling tradeoffs. The $\mathcal{O}(N^3)$ computational scaling inherent to exact GP inference necessitated aggressive sub-sampling of the high-cadence Kepler observations. This sub-sampling intrinsically suppressed high-frequency variability, limiting kernel discriminability and reducing the effective information available to the optimizer. Sparse GP approximations—such as inducing point methods or Sparse Variational Gaussian Processes (SVGP)—would enable the utilization of full datasets and likely restore discriminability between kernel classes on the rotating-star problem.

Block gap extrapolation limits. While the GP maintained stable reconstruction within block gaps, the posterior variance in the gap centers reflects genuine epistemic uncertainty. Extended gaps exceeding multiple characteristic length-scales necessarily revert to the prior, and no inference procedure can recover information absent from the training data in such settings.

9 Conclusion

This study demonstrates that Gaussian Process regression provides a robust and probabilistically calibrated framework for imputing missing observations in quiet stellar light curves, while also revealing principled limitations when applied to highly nonstationary rotating-star variability. Overall, these experiments suggest that Gaussian Process methods provide highly effective and probabilistically calibrated imputation for relatively stationary stellar variability regimes, while globally stationary covariance assumptions become increasingly inadequate under evolving variance structure.

For the quiet-star dataset (KIC 007596240), GP regression successfully reconstructed missing flux values with mean RMSE of 5.0477 ± 0.2525 across 20 Monte Carlo splits, outperforming both linear interpolation (6.2827 ± 0.3437) and cubic spline interpolation (8.6941 ± 1.1385). Predictive uncertainty calibration remained close to ideal Gaussian expectations (1σ : $65.12\% \pm 4.24\%$; 2σ : $96.07\% \pm 1.63\%$), confirming that the GP posterior variance constitutes a scientifically reliable imputation confidence metric. Block gap reconstruction experiments further demonstrated that the GP maintains stable and physically plausible predictions under extended observational interruptions, with uncertainty growing gracefully rather than diverging.

For the rotating-star dataset (KIC 007609553), the quasi-periodic kernel did not achieve statistically meaningful improvements over the standard RBF kernel. This outcome is directly attributable to the strong nonstationarity of the rotating-star variability: rolling variance analysis confirmed that the local second-order statistics evolve substantially across the observational baseline, violating the stationarity assumptions shared by both kernel classes. Kernel sophistication alone cannot compensate for this structural model mismatch. Furthermore, the CP-GP experiment revealed that a piecewise-stationary approach also fails to yield substantial predictive gains, suggesting that evolving stellar variance may require more flexible, continuously nonstationary probabilistic models.

The experiments therefore highlight a scientifically meaningful boundary of the stationary GP framework, motivating future adoption of fully nonstationary covariance models, heteroscedastic architectures, and sparse inference methods to extend probabilistic stellar variability modeling to active rotating stars at full cadence.

References

- [1] Suzanne Aigrain and Daniel Foreman-Mackey. *Gaussian Process regression for astronomical time-series*. 2022. arXiv: 2209.08940 [astro-ph.IM]. URL: <https://arxiv.org/abs/2209.08940>.
- [2] Ruth Angus et al. “Inferring probabilistic stellar rotation periods using Gaussian processes”. In: *Monthly Notices of the Royal Astronomical Society* 474.2 (2017), pp. 2094–2108. ISSN: 1365-2966. DOI: 10.1093/mnras/stx2109. URL: <http://dx.doi.org/10.1093/mnras/stx2109>.
- [3] William J. Borucki et al. “Kepler Planet-Detection Mission: Introduction and First Results”. In: *Science* 327.5968 (2010), pp. 977–980. DOI: 10.1126/science.1185402. eprint: <https://www.science.org/doi/pdf/10.1126/science.1185402>. URL: <https://www.science.org/doi/abs/10.1126/science.1185402>.
- [4] Eric Feigelson, G. Babu, and Gabriel Caceres. “Autoregressive Times Series Methods for Time Domain Astronomy”. In: *Frontiers in Physics* 6 (Aug. 2018), p. 80. DOI: 10.3389/fphy.2018.00080.
- [5] Daniel Foreman-Mackey et al. “Fast and Scalable Gaussian Process Modeling with Applications to Astronomical Time Series”. In: *The Astronomical Journal* 154.6 (Nov. 2017), p. 220. ISSN: 1538-3881. DOI: 10.3847/1538-3881/aa9332. URL: <http://dx.doi.org/10.3847/1538-3881/aa9332>.
- [6] B A Nicholson and S Aigrain. “Quasi-periodic Gaussian processes for stellar activity: From physical to kernel parameters”. In: *Monthly Notices of the Royal Astronomical Society* 515.4 (2022), pp. 5251–5266. ISSN: 1365-2966. DOI: 10.1093/mnras/stac2097. URL: <http://dx.doi.org/10.1093/mnras/stac2097>.
- [7] Niamh K. O’Sullivan and Suzanne Aigrain. *Modelling stochastic and quasi-periodic behaviour in stellar time-series: Gaussian process regression versus power-spectrum fitting*. 2024. arXiv: 2404.11662 [astro-ph.SR]. URL: <https://arxiv.org/abs/2404.11662>.
- [8] Christopher Paciorek and Mark Schervish. “Nonstationary Covariance Functions for Gaussian Process Regression”. In: *Advances in Neural Information Processing Systems* 16 (Mar. 2004).
- [9] Carl Edward Rasmussen and Christopher K. I. Williams. *Gaussian Processes for Machine Learning*. MIT Press, 2006.